

Go execution tracer

Gophercon Brazil 2017

André Carvalho

\$ whoami

- Developer @ globo.com / tsuru
- Interested in all things performance, tracing and systems programming
- @andresantostc
- <https://andrestc.com>



Agenda

- The tool basics
 - What is the Go execution tracer
 - How to collect trace data
 - How to analyze trace data
- Using the tracer
 - to improve working code
 - to investigate a strange behavior

Go execution tracer

- Gives insight into the execution of a Go program
 - What are my goroutines doing when not in CPU?
- Instruments the go runtime
 - Captures events in nanosecond precision
 - Data is not aggregated/sampled!
- Available since go 1.5

Go execution tracer

- Events
 - Goroutines creation/start/end/block/unblock
 - Network
 - Syscalls
 - Memory allocation
 - Garbage collection

Collecting traces

- Three ways to collect traces
 - `trace.Start / trace.Stop`
 - `go test -trace=trace.out`
 - `debug/pprof/trace` handler

Collecting traces

- Writes the tracing output to os.Stderr
- go run main.go 2> trace.out
- go tool trace trace.out

```
1  package main
2
3  import (
4      "os"
5      "runtime/trace"
6  )
7
8  func main() {
9      trace.Start(os.Stderr)
10     defer trace.Stop()
11     // create new channel of type int
12     ch := make(chan int)
13
14     // start new anonymous goroutine
15     go func() {
16         // send 42 to channel
17         ch ← 42
18     }()
19     // read from channel
20     ←ch
21 }
```

Trace

[View trace](#)

[Goroutine analysis](#)


[Network blocking profile](#)

[Synchronization blocking profile](#)

[Syscall blocking profile](#)

[Scheduler latency profile](#)

View Trace



1. Timeline

2. Heap usage

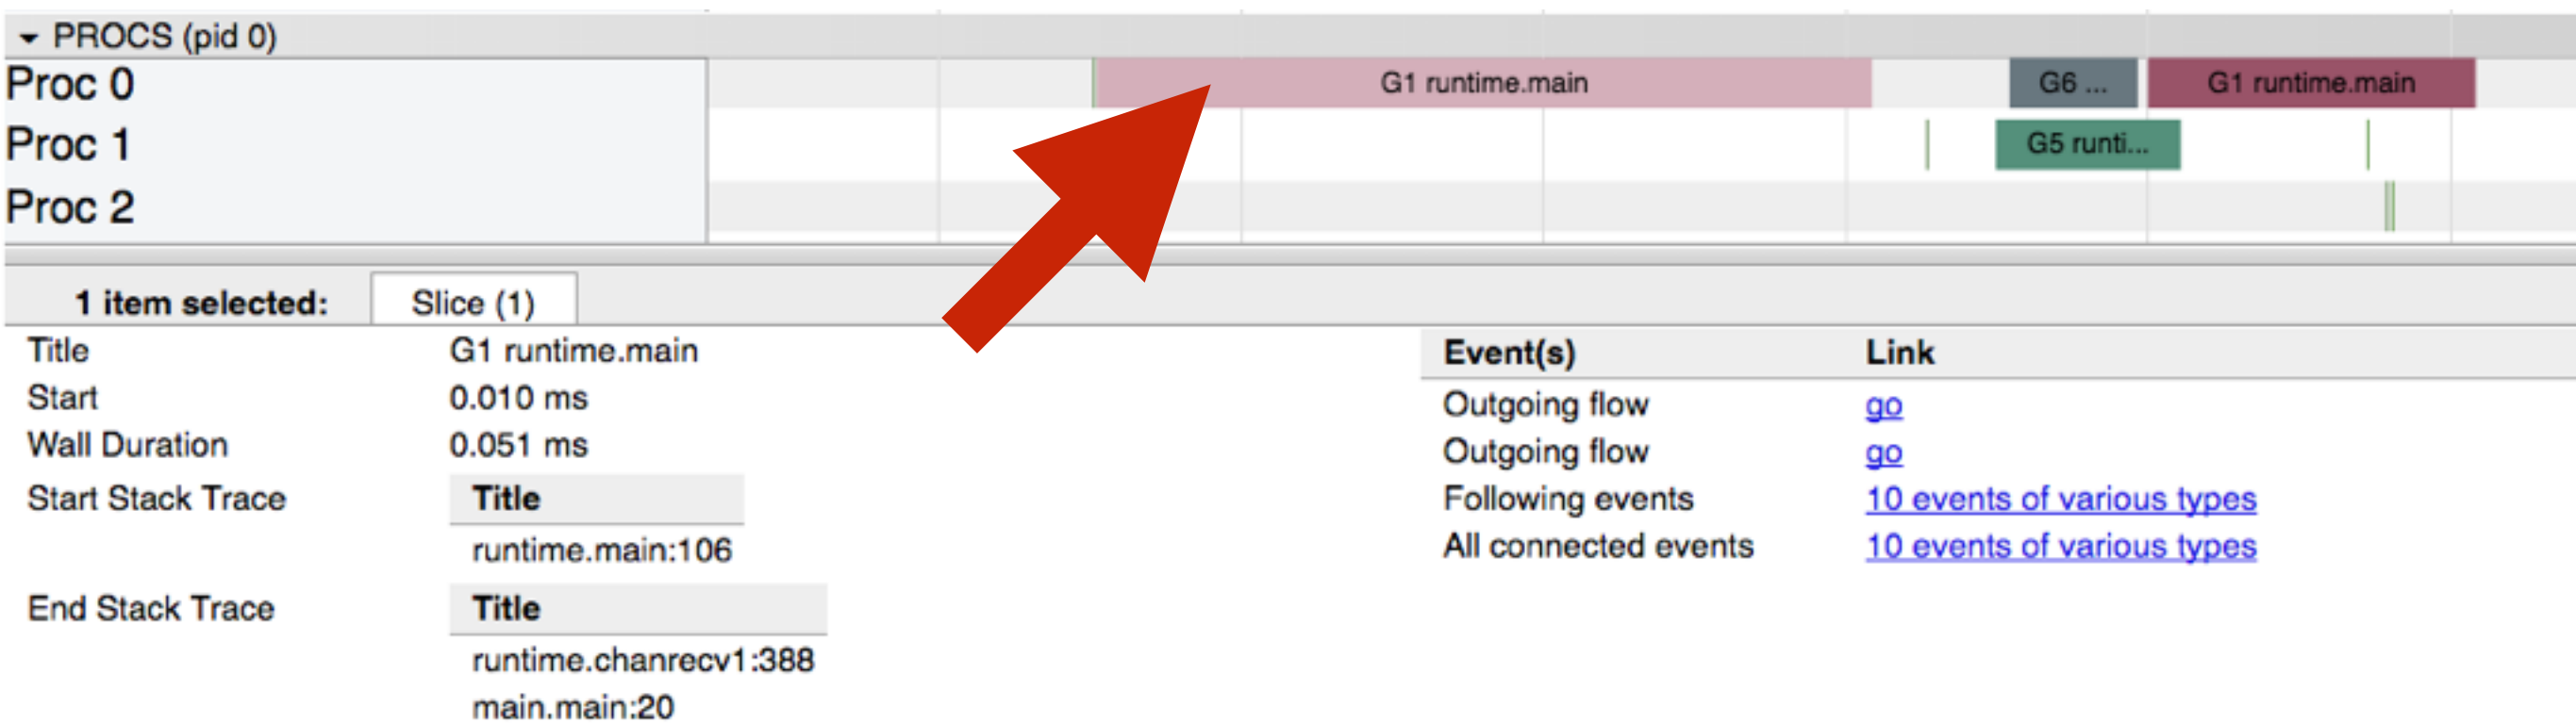
3. goroutines

4. OS threads

5. Virtual Processors

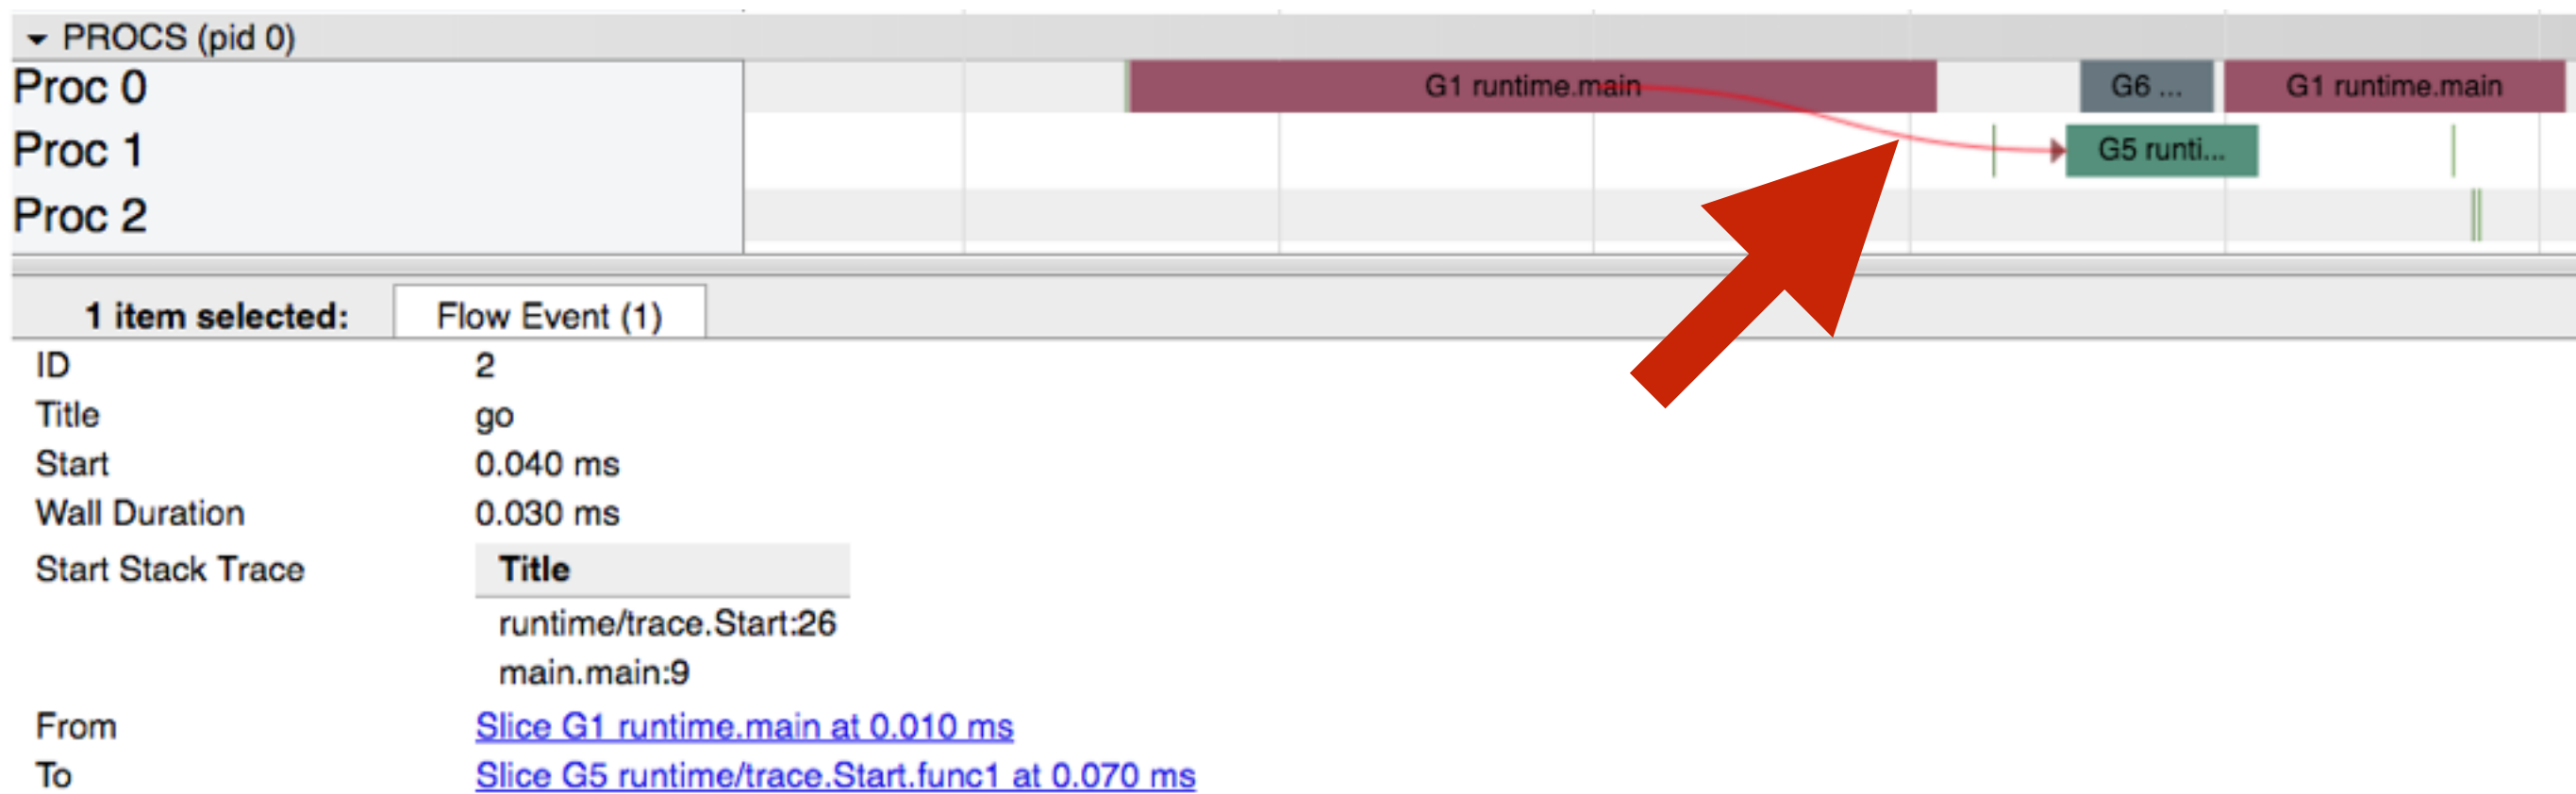
6. goroutine + events

View Trace



Selecting a goroutine

View Trace



Selecting an event

Tracing web applications

```
1  package main
2
3  import (
4      "net/http"
5      _ "net/http/pprof"
6  )
7
8  func main() {
9      http.Handle("/hello", http.HandlerFunc(helloHandler))
10
11     http.ListenAndServe(":8080", http.DefaultServeMux)
12 }
13
14 func helloHandler(w http.ResponseWriter, r *http.Request) {
15     w.Write([]byte("hello world!"))
16 }
```


<https://github.com/andrestc/go-tracing/tree/master/07-web>

Using the tracer to
improve working code

Mandelbrot

CPU intensive calculations to
figure out each pixel's color

[https://github.com/campoy/
mandelbrot/](https://github.com/campoy/mandelbrot/)



Pixel

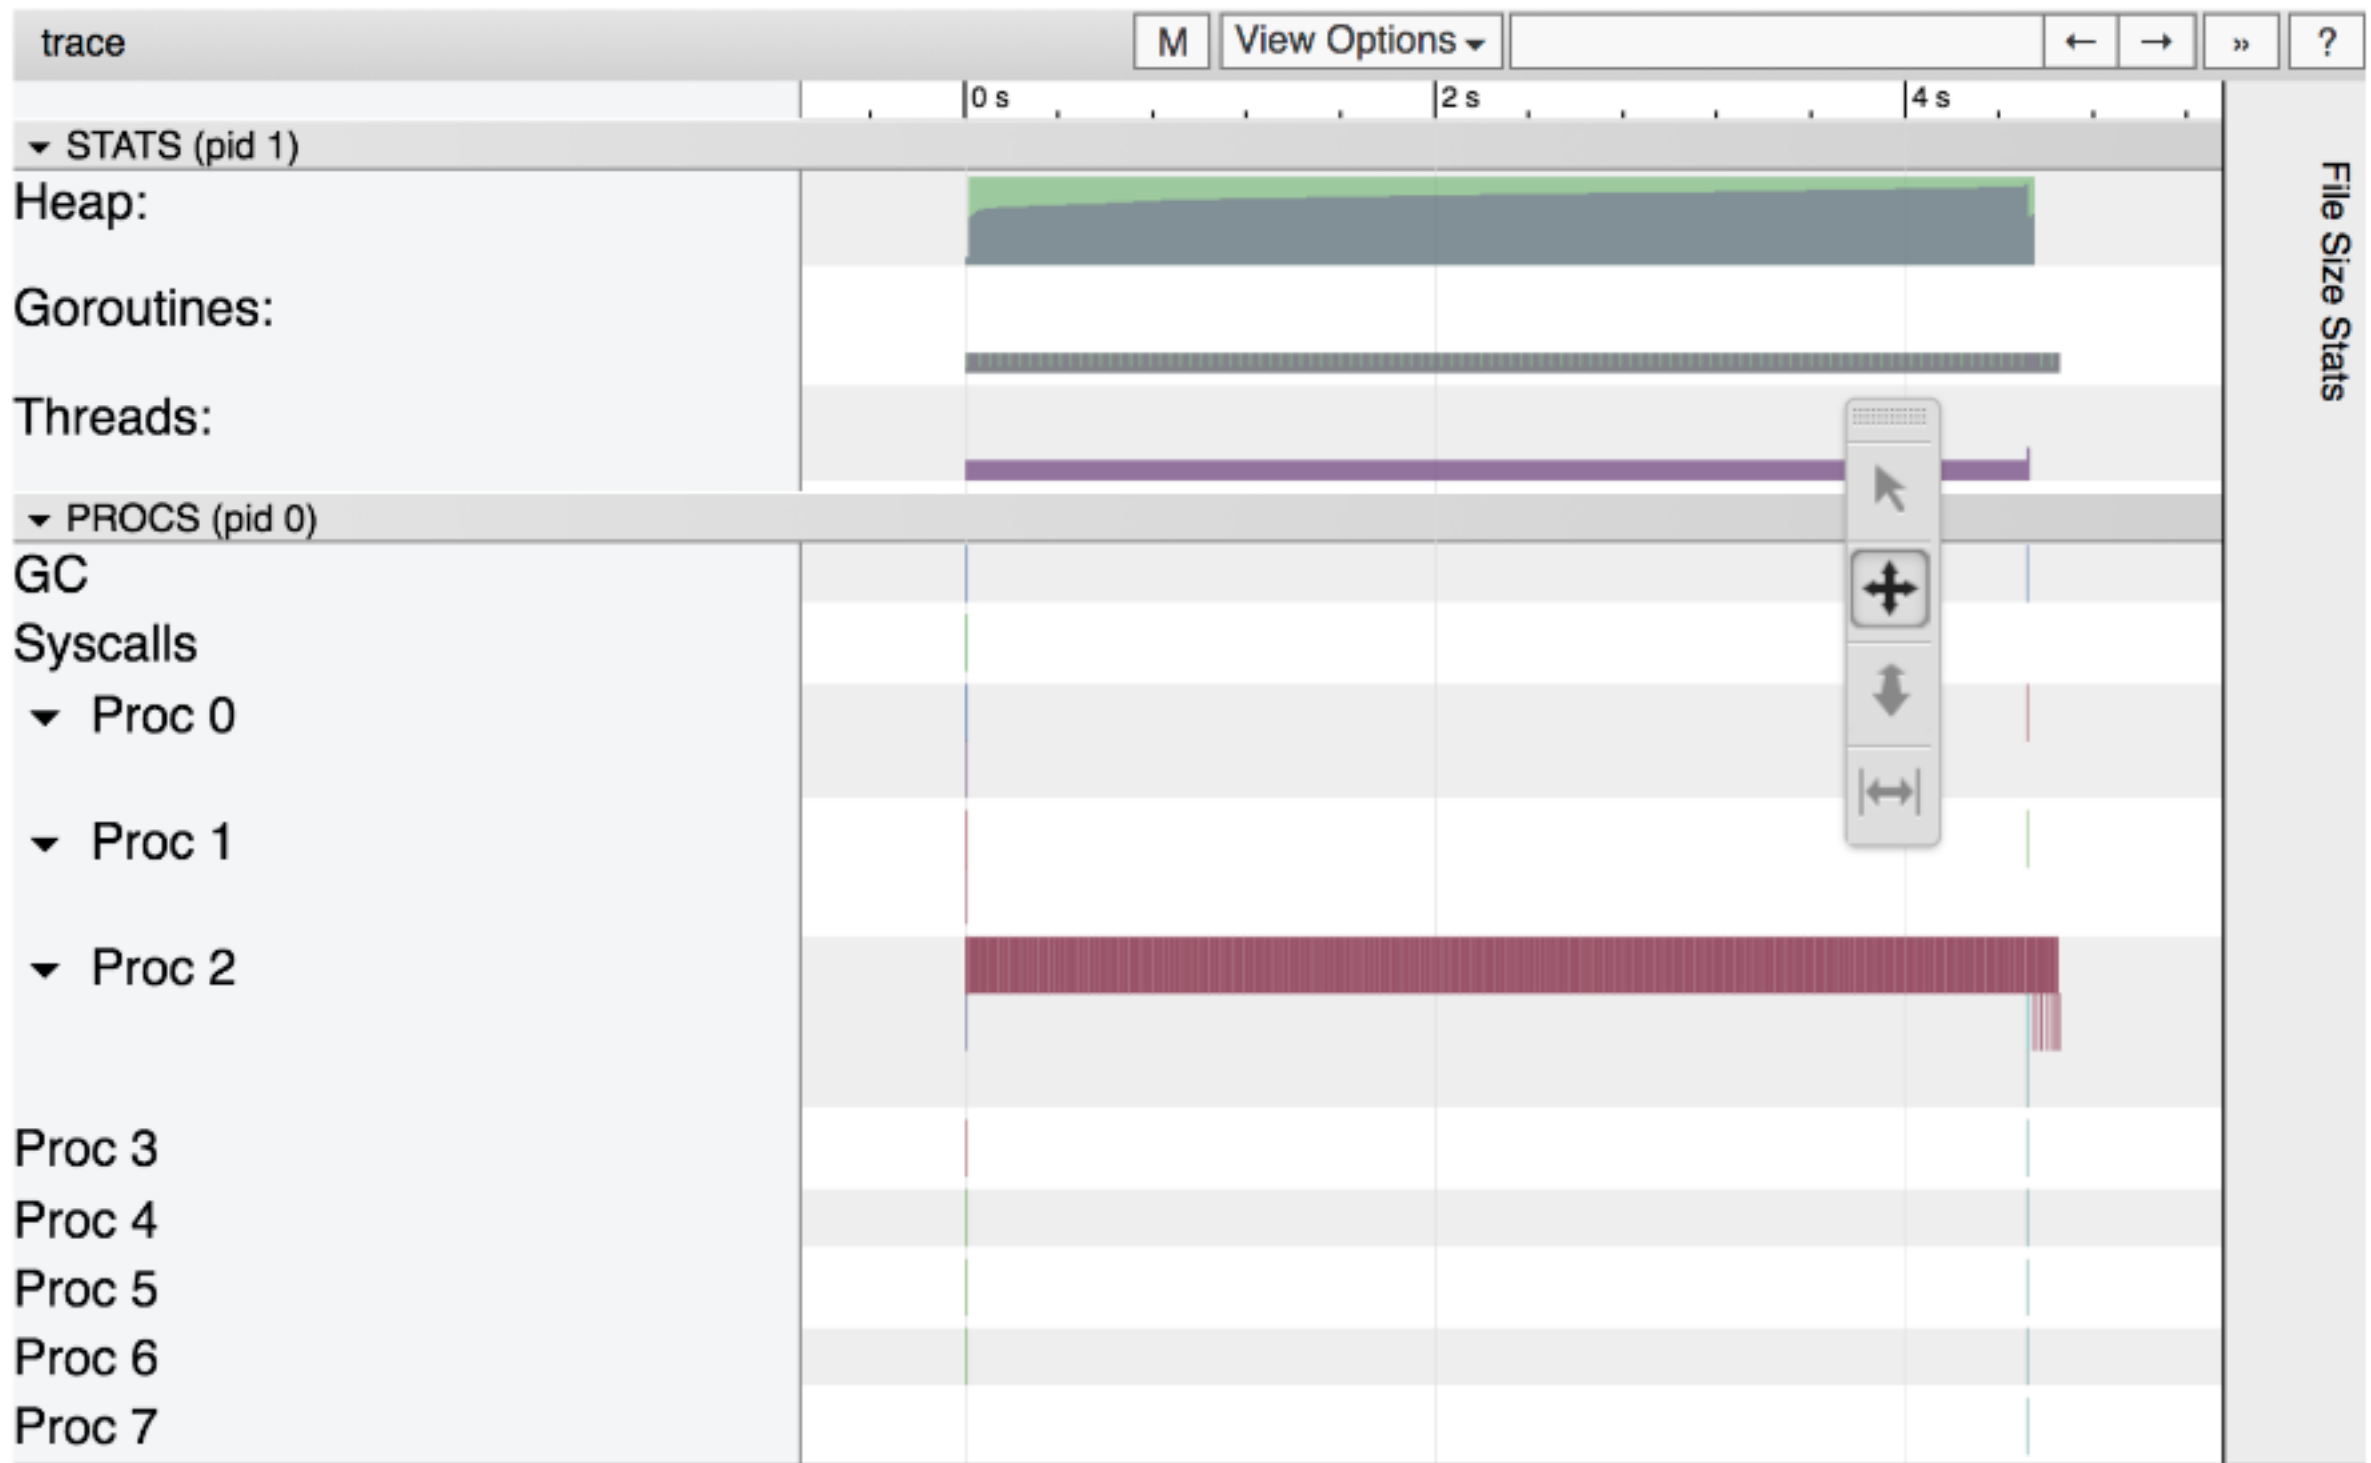
```
func pixel(i, j, width, height int) color.Color {  
    const complexity = 4  
  
    xi := norm(i, width, -1.0, 2)  
    yi := norm(j, height, -1, 1)  
  
    const maxI = 1000  
    x, y := 0., 0.  
  
    for i := 0; (x*x+y*y < complexity) && i < maxI; i++ {  
        x, y = x*x-y*y+xi, 2*x*y+yi  
    }  
  
    return color.Gray{uint8(x)}  
}
```

Sequential

```
func createSeq(width, height int) image.Image {  
    m := image.NewGray(image.Rect(0, 0, width, height))  
    for i := 0; i < width; i++ {  
        for j := 0; j < height; j++ {  
            m.Set(i, j, pixel(i, j, width, height))  
        }  
    }  
    return m  
}
```

```
$ time ./mandelbrot  
./mandelbrot 4.44s user 0.01s system 99% cpu 4.466 total
```


Sequential

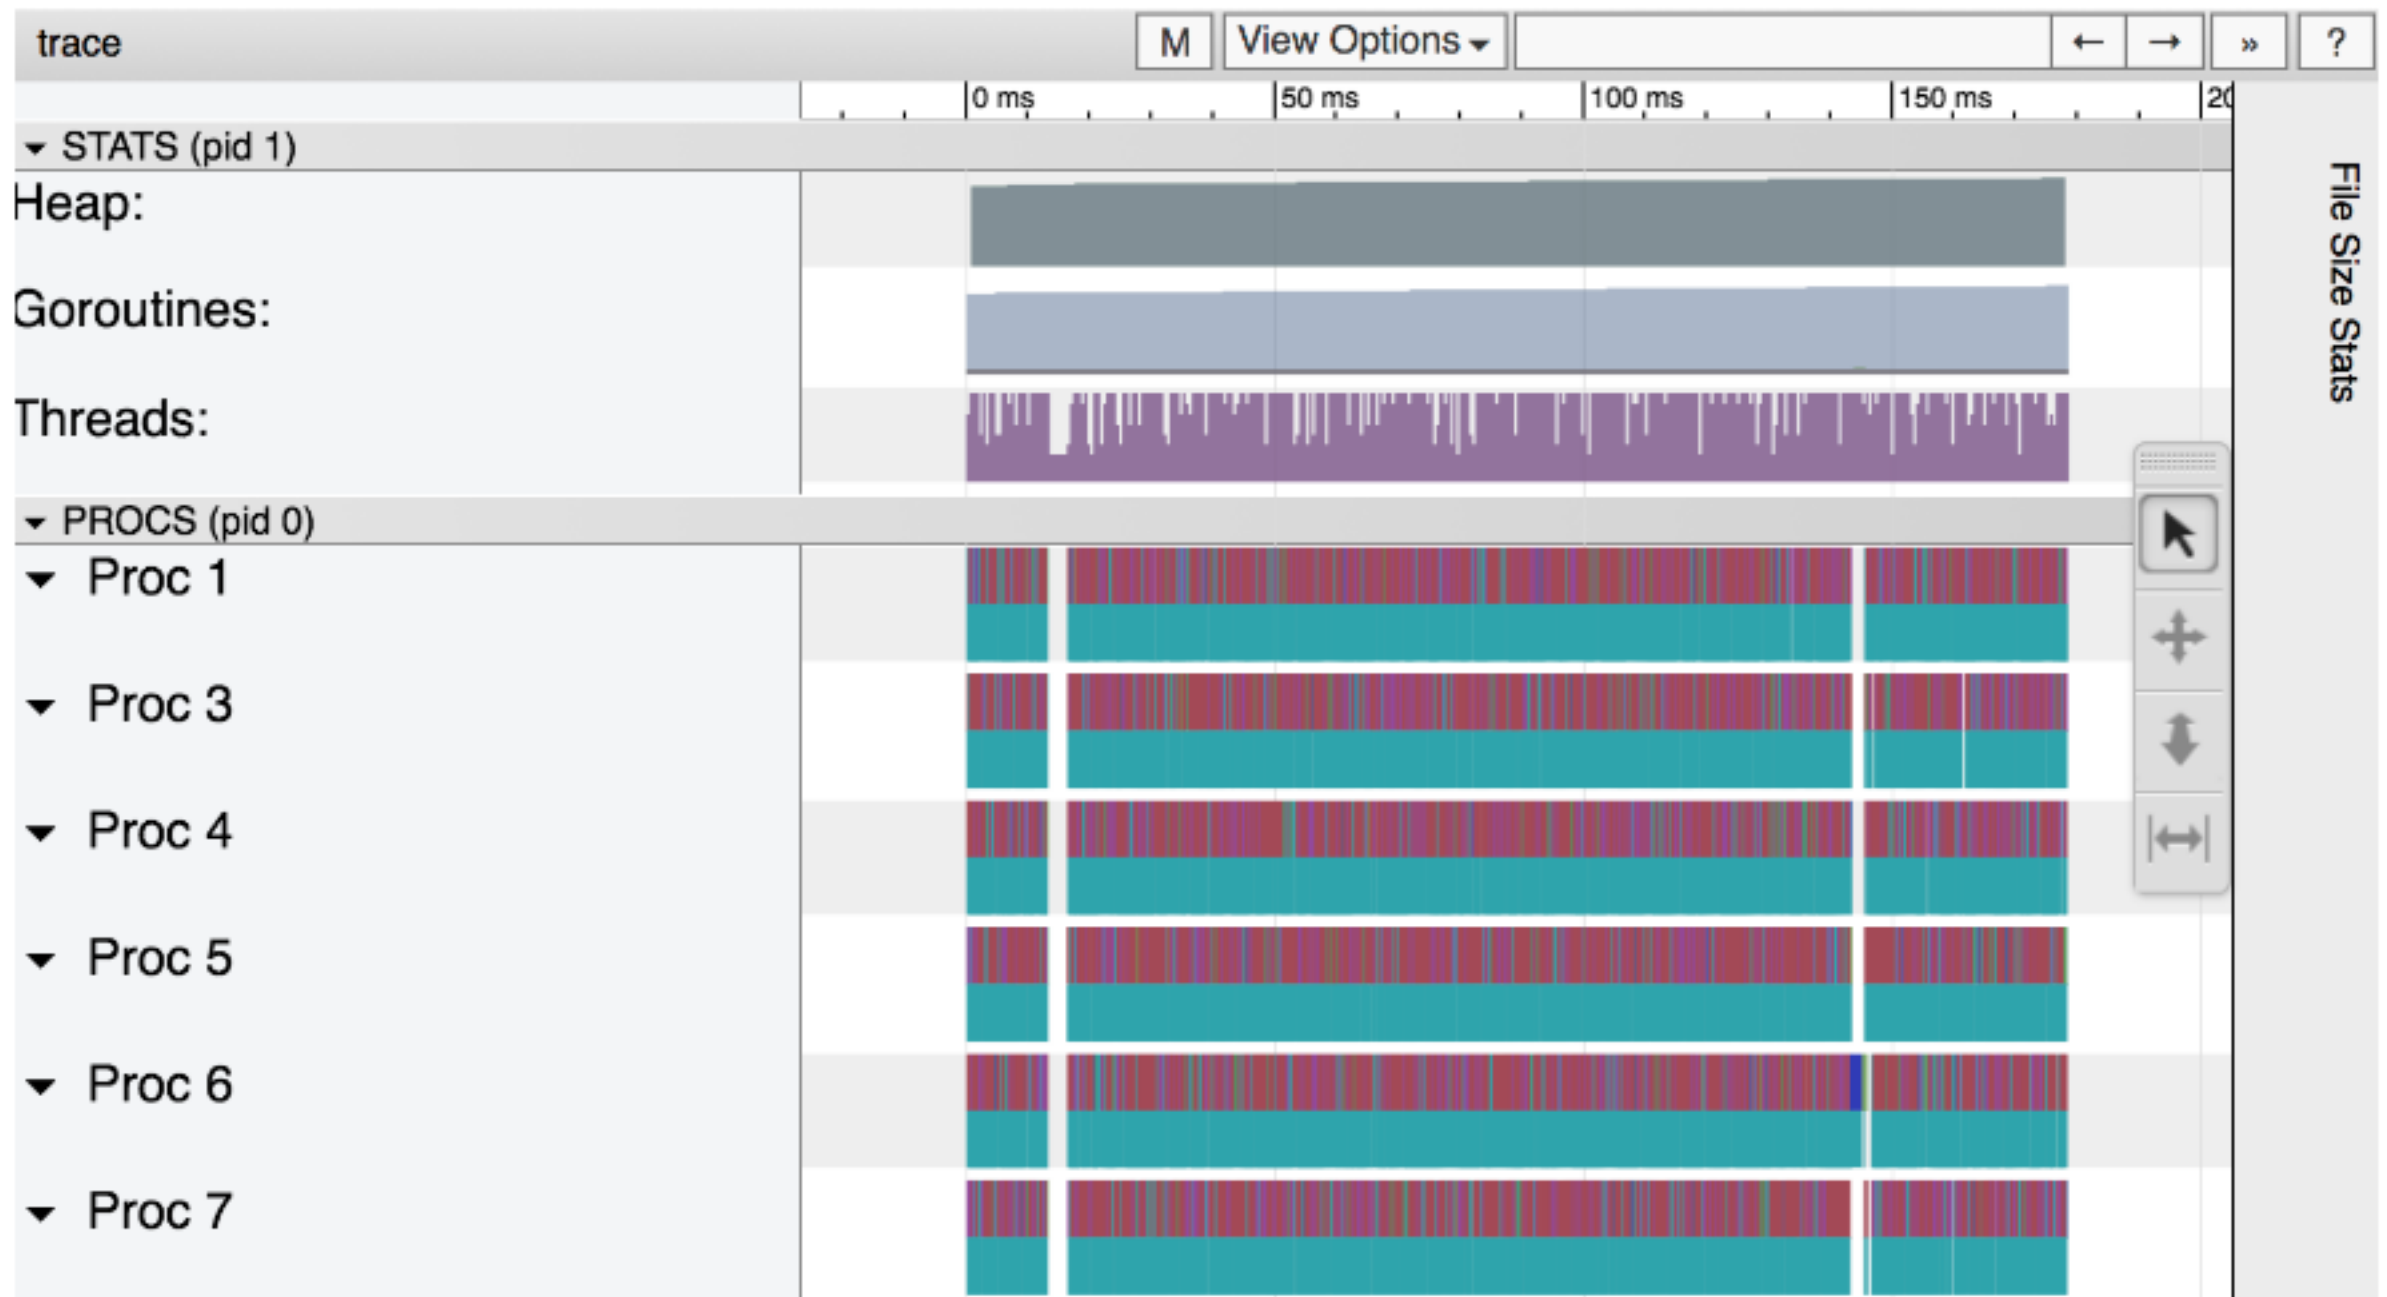


One goroutine per pixel

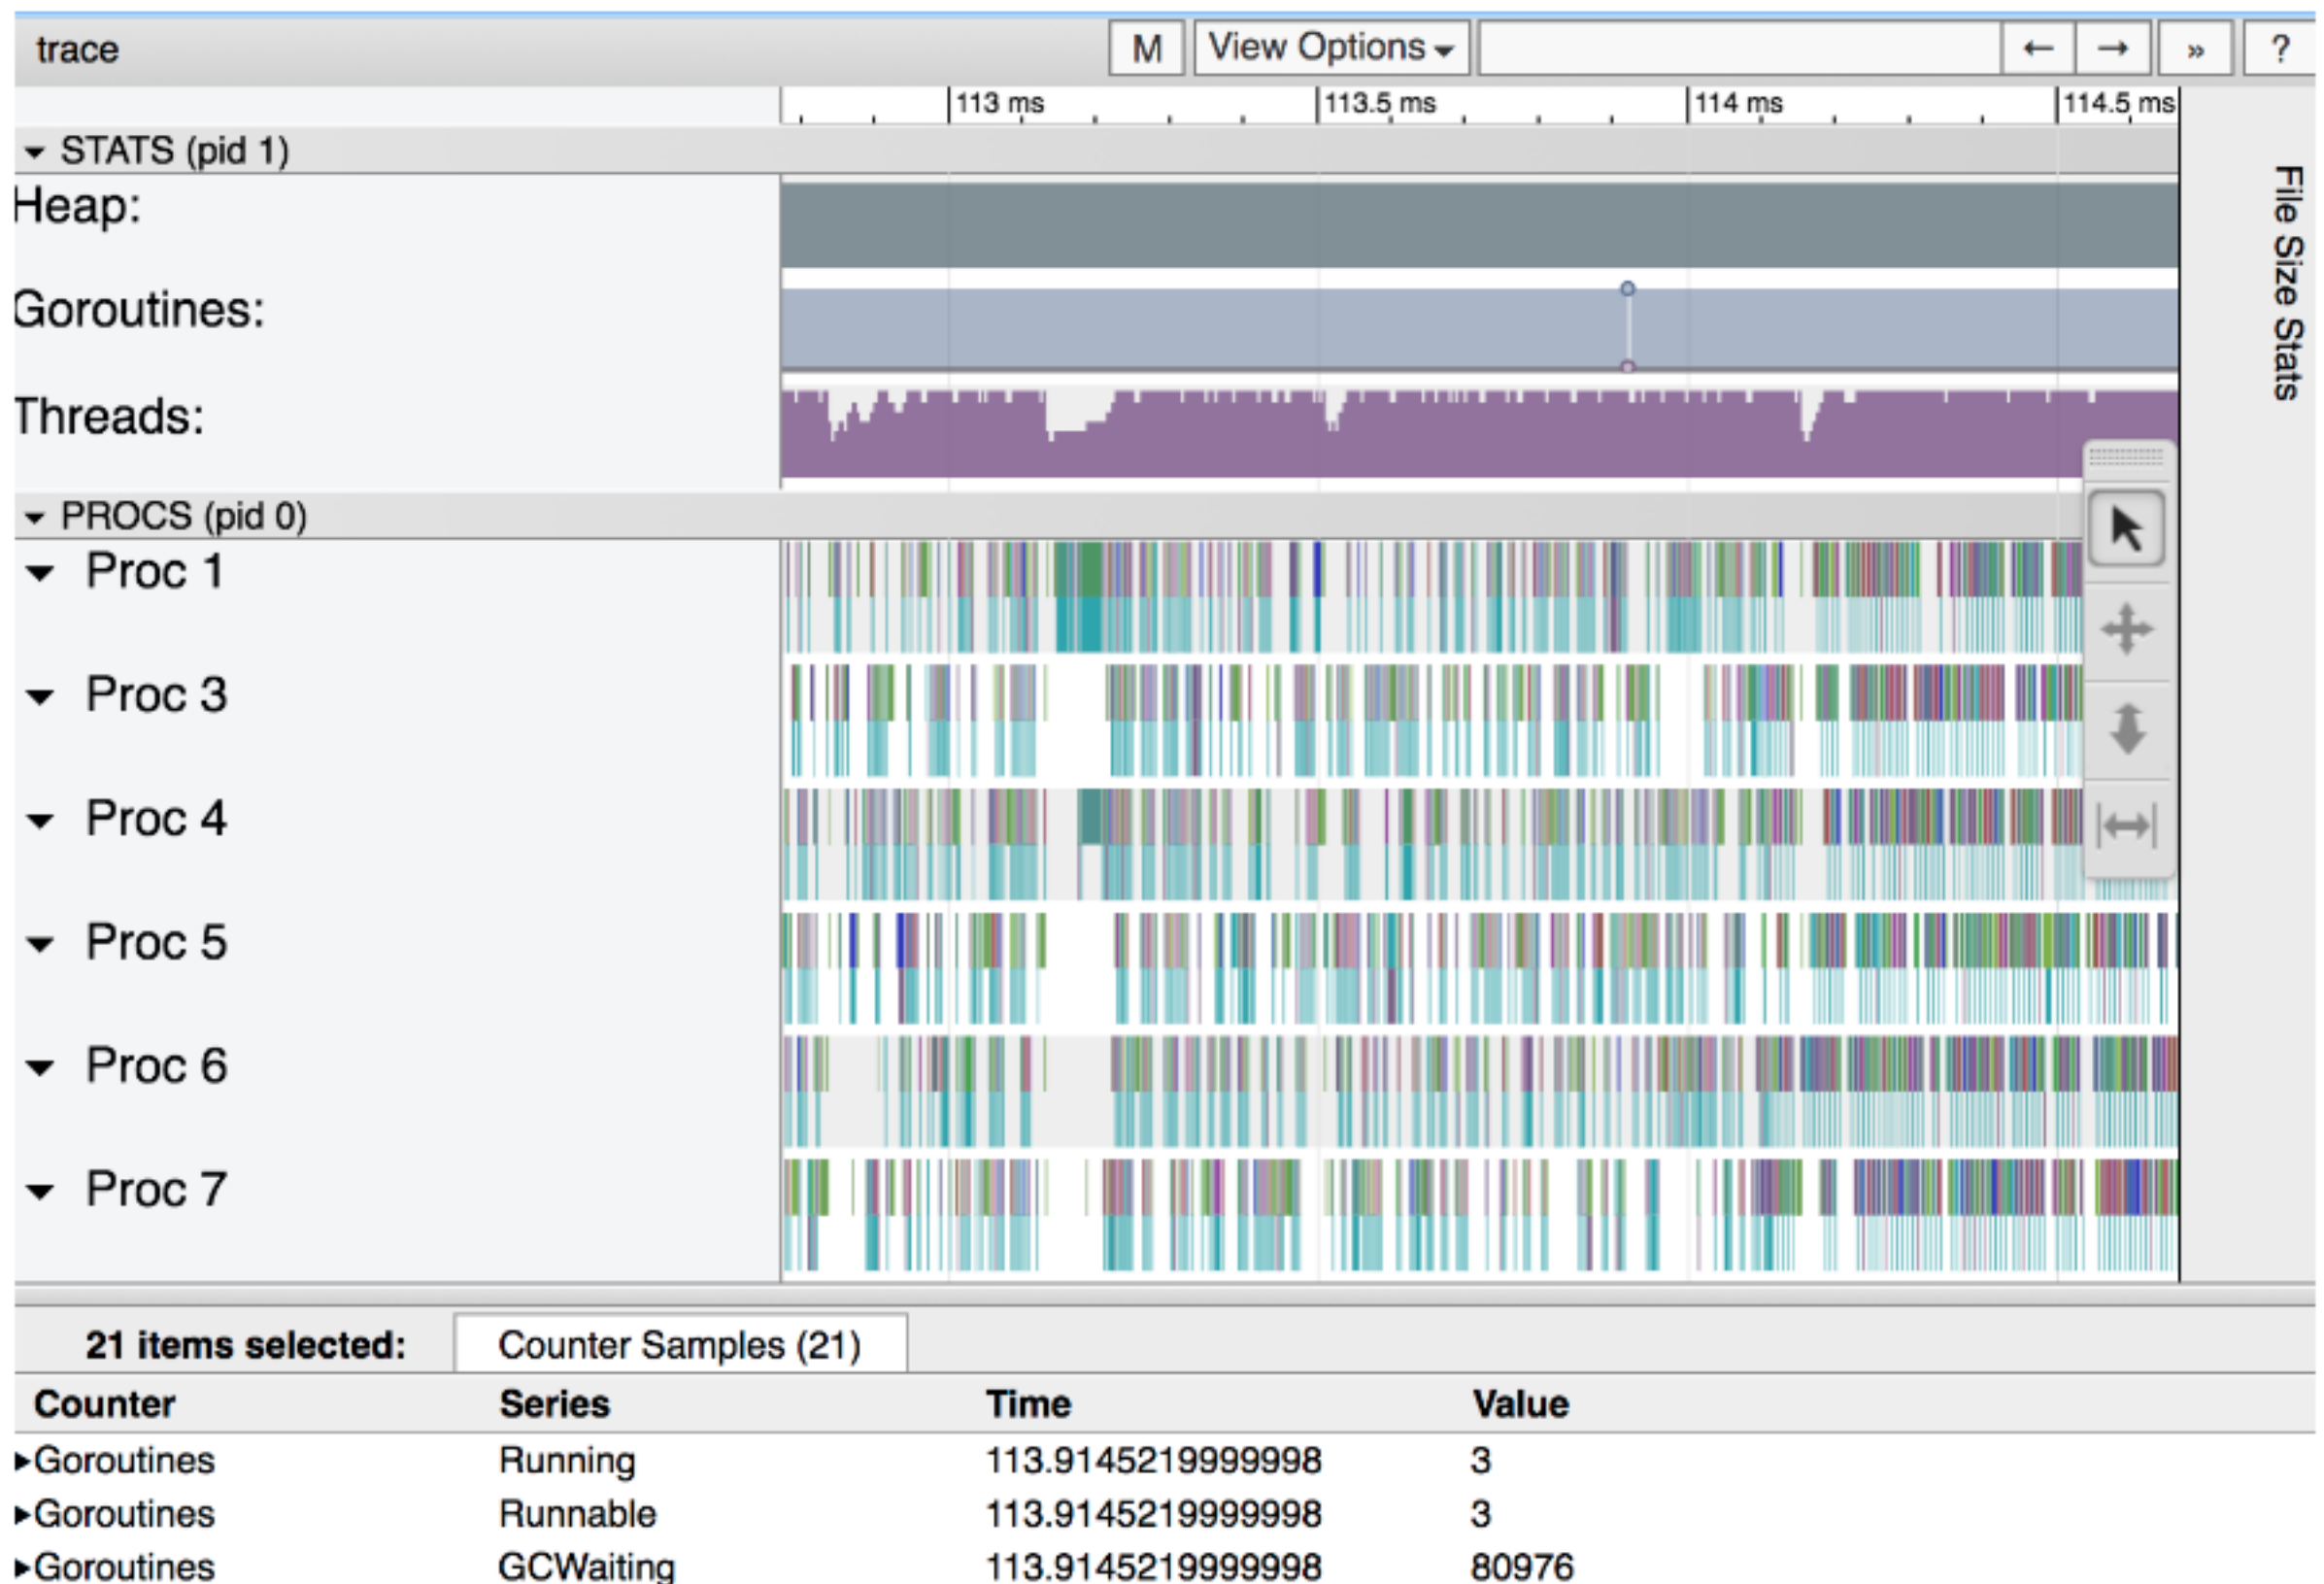
```
func createPixel(width, height int) image.Image {  
    m := image.NewGray(image.Rect(0, 0, width, height))  
    var w sync.WaitGroup  
    w.Add(width * height)  
    for i := 0; i < width; i++ {  
        for j := 0; j < height; j++ {  
            go func(i, j int) {  
                m.Set(i, j, pixel(i, j, width, height))  
                w.Done()  
            }(i, j)  
        }  
    }  
    w.Wait()  
    return m  
}
```

```
$ time ./mandelbrot  
./mandelbrot 13.70s user 1.53s system 471% cpu 3.234 total
```

One goroutine per pixel



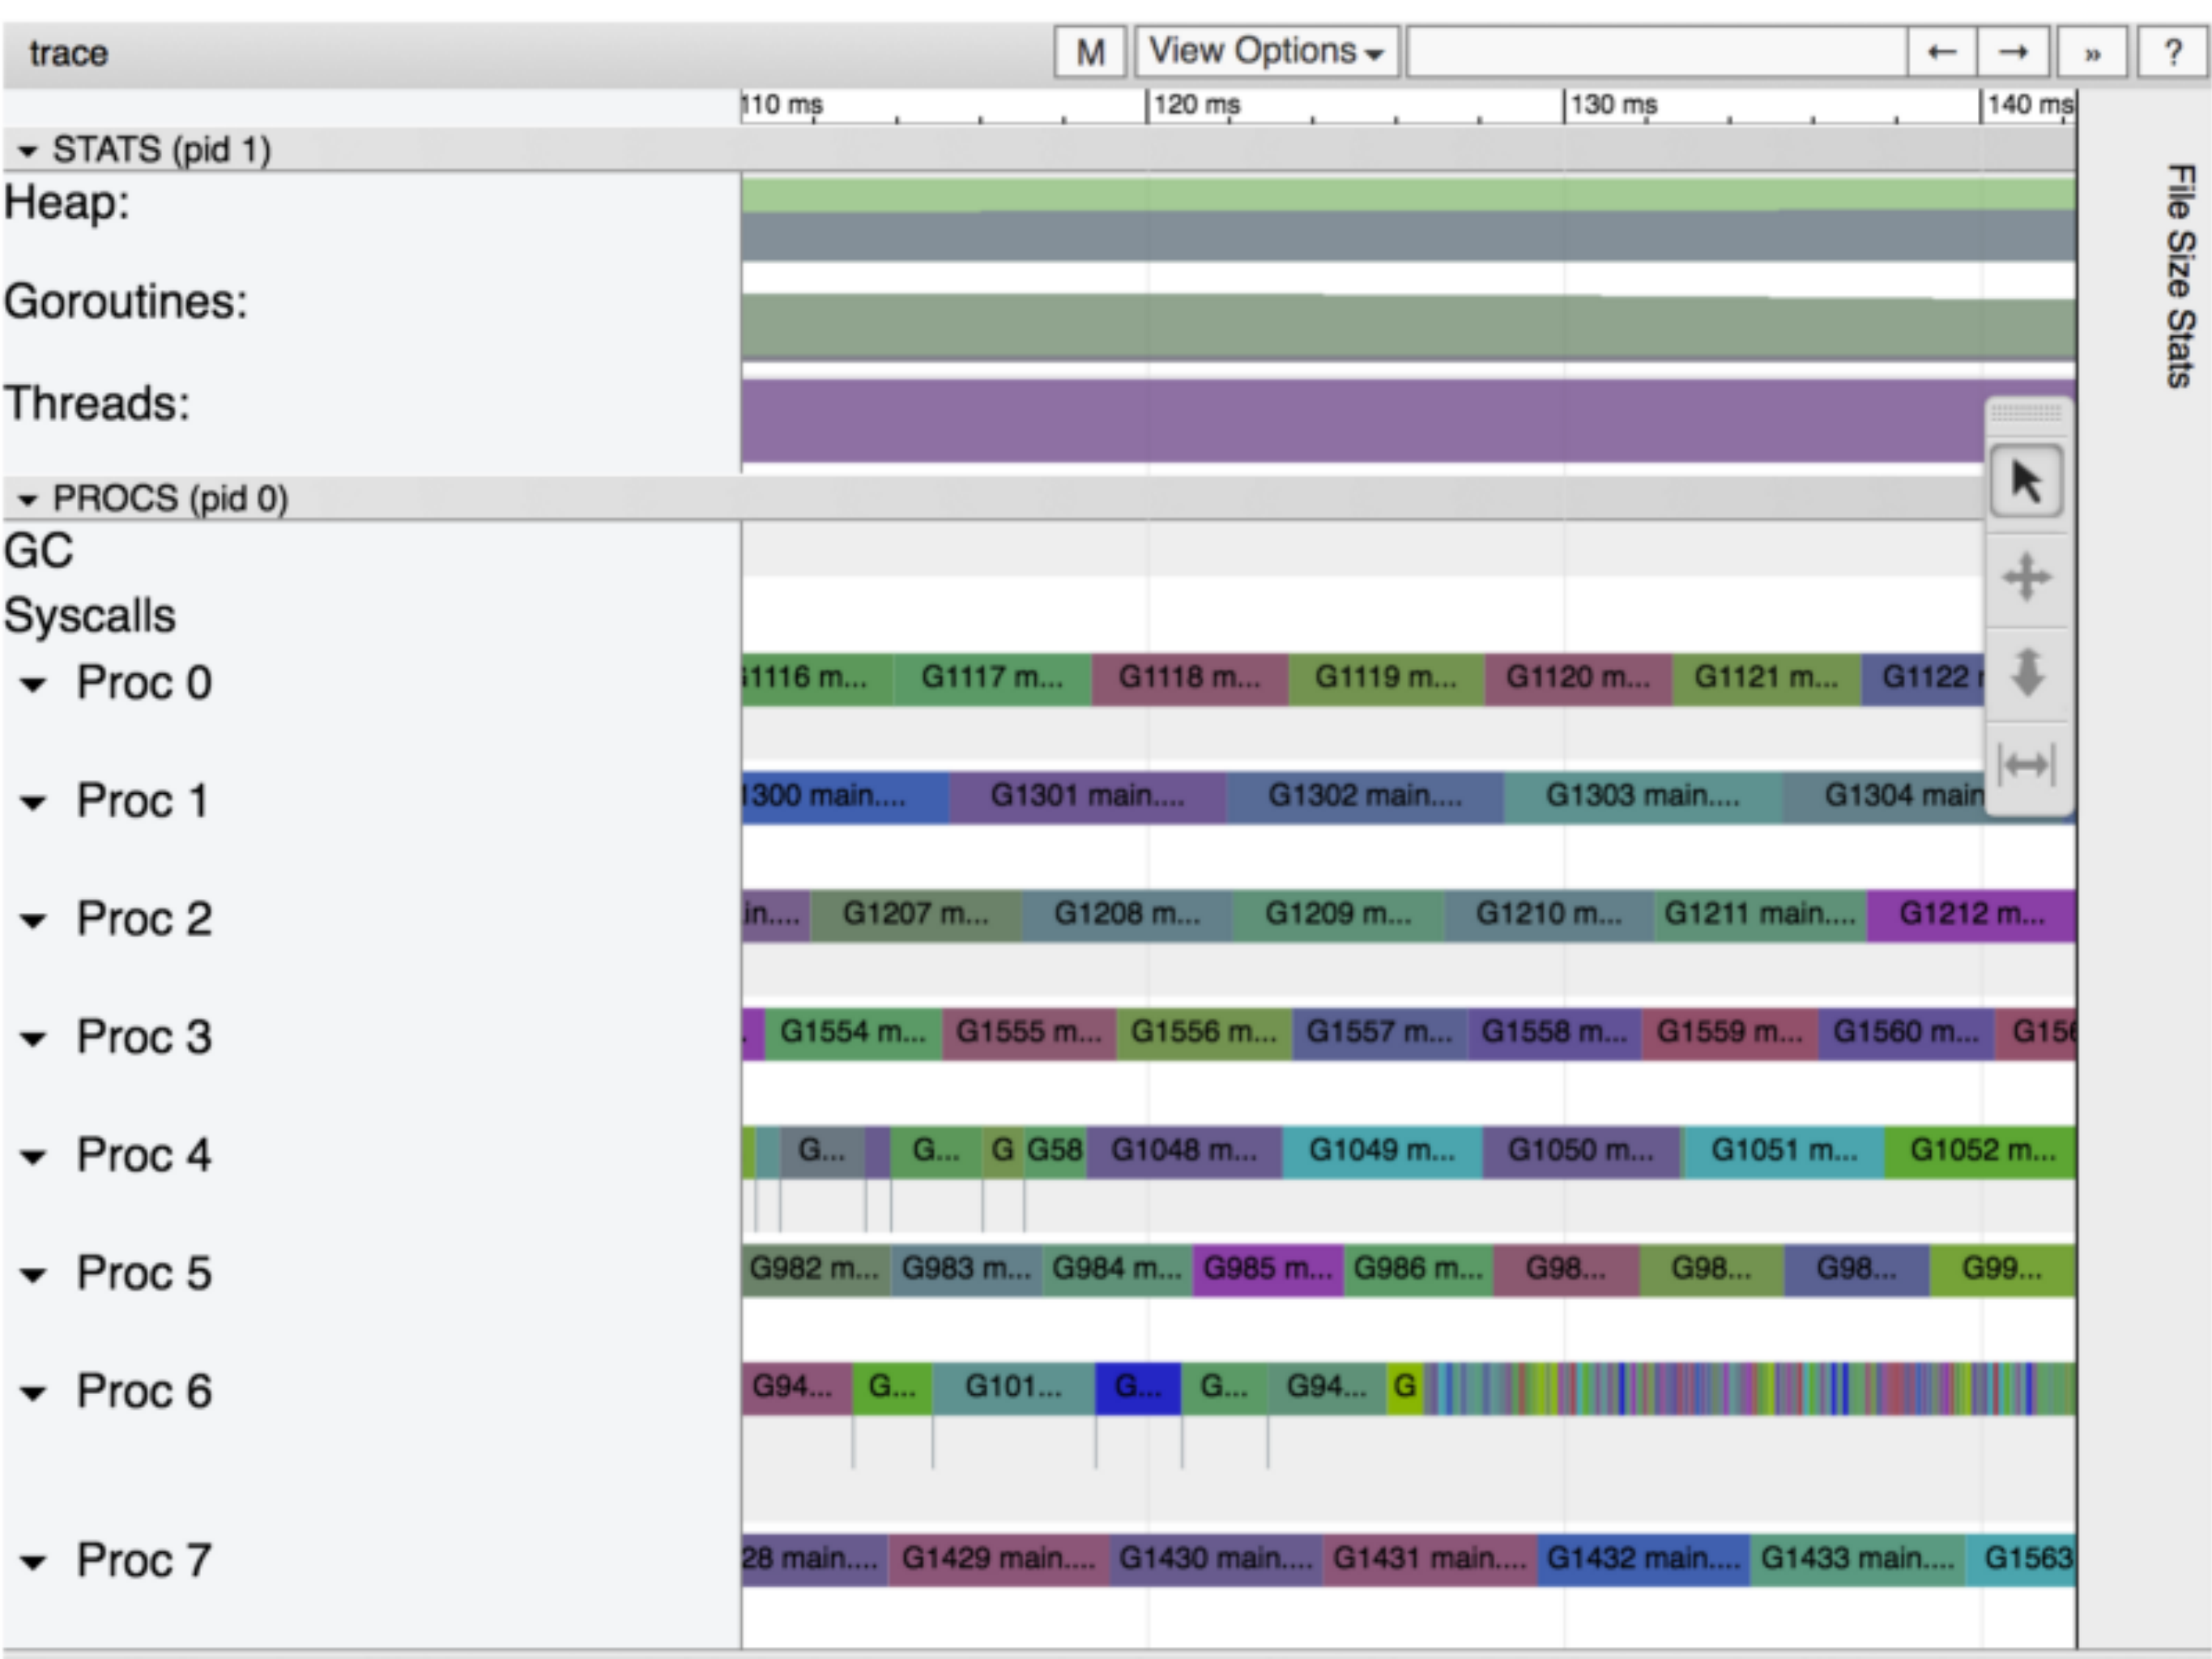
One goroutine per pixel



One goroutine per column

```
func createCol(width, height int) image.Image {  
    m := image.NewGray(image.Rect(0, 0, width, height))  
    var w sync.WaitGroup  
    w.Add(width)  
    for i := 0; i < width; i++ {  
        go func(i int) {  
            for j := 0; j < height; j++ {  
                m.Set(i, j, pixel(i, j, width, height))  
            }  
            w.Done()  
        }(i)  
    }  
    w.Wait()  
    return m  
}
```

```
$ time ./mandelbrot  
./mandelbrot 4.90s user 0.02s system 499% cpu 0.985 total
```

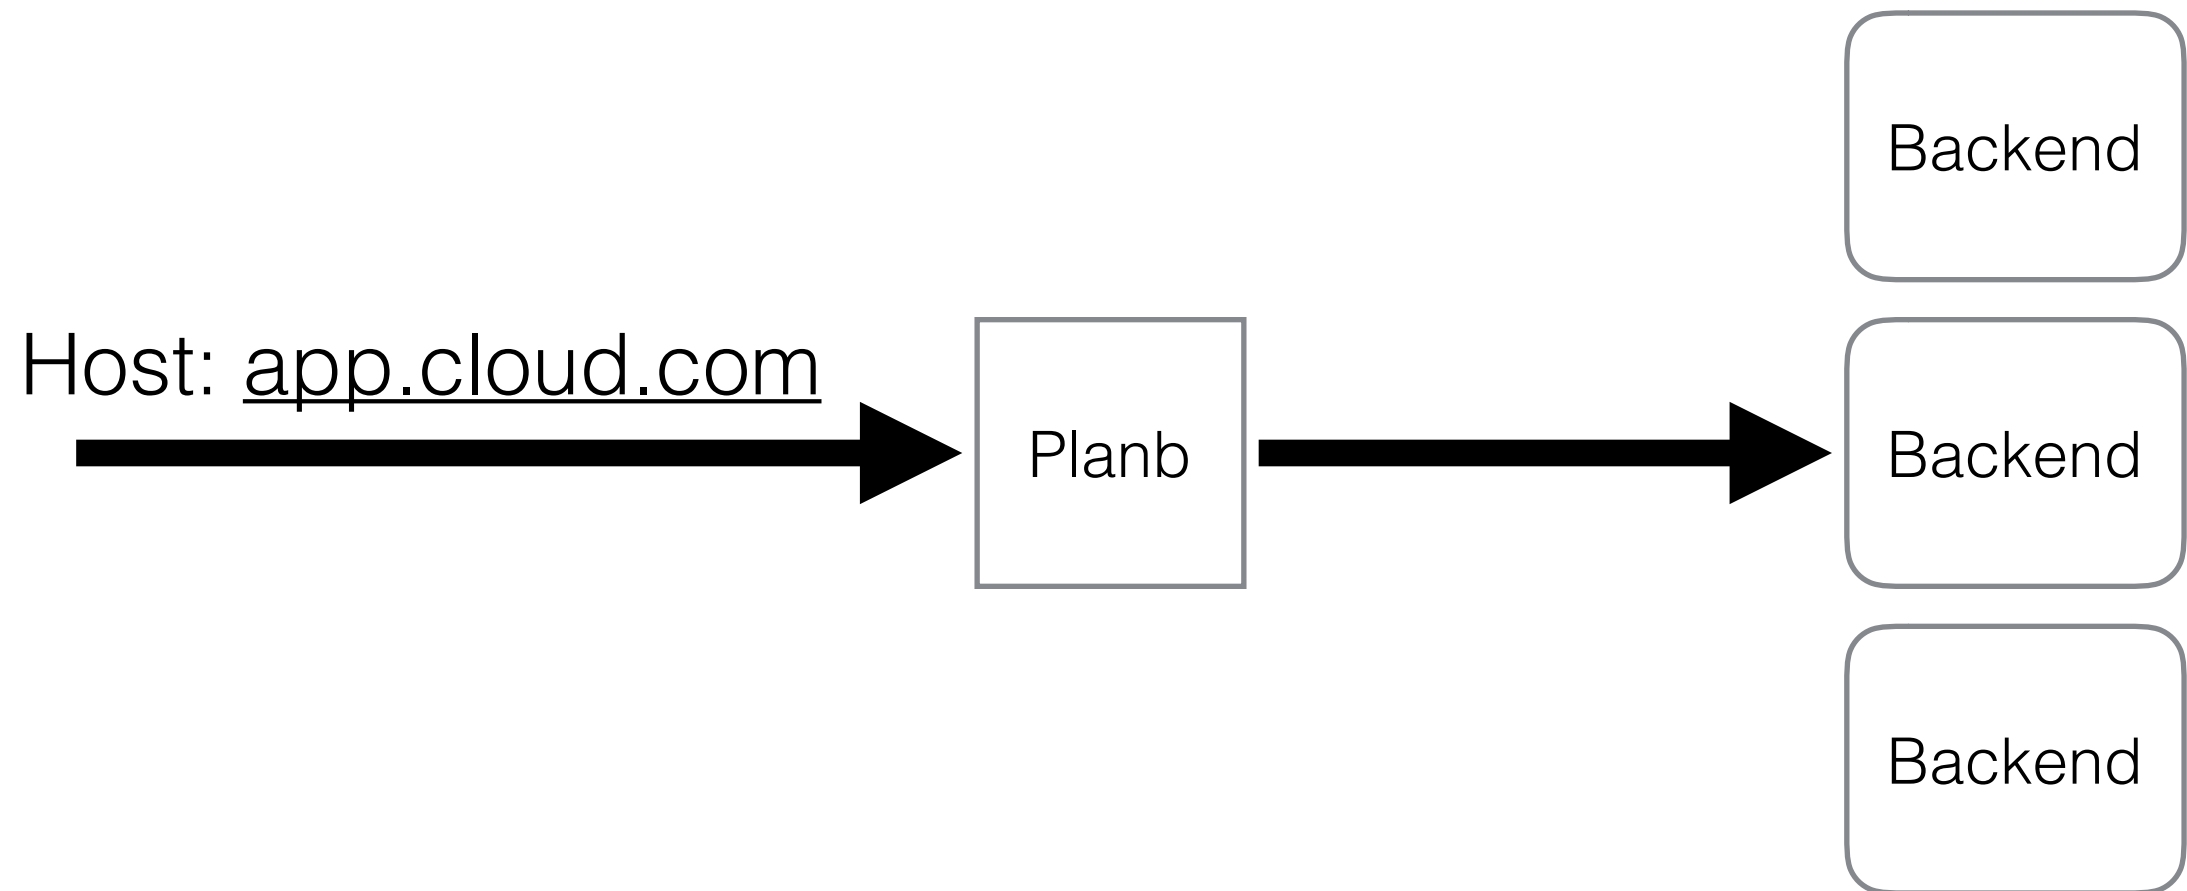


Debugging a strange
test behavior

Planb

- github.com/tsuru/planb
 - reverse proxy based of [hipache](#)
- One of tsuru's key components
- Used extensively on globo.com

Planb



```
$ go test -check.v -check.f TestRoundTripStressWithTimeout
PASS: reverseproxy_test.go:763: S.TestRoundTripStressWithTimeoutBackend 0.012s
PASS: reverseproxy_test.go:763: S.TestRoundTripStressWithTimeoutBackend 0.011s
OK: 2 passed
PASS
ok      github.com/tsuru/planb/reverseproxy    0.044s
```

```
$ go test -check.v -check.f TestRoundTripStressWithTimeout
PASS: reverseproxy_test.go:763: S.TestRoundTripStressWithTimeoutBackend 60.013s
PASS: reverseproxy_test.go:763: S.TestRoundTripStressWithTimeoutBackend 0.014s
OK: 2 passed
PASS
ok      github.com/tsuru/planb/reverseproxy    60.045s
```

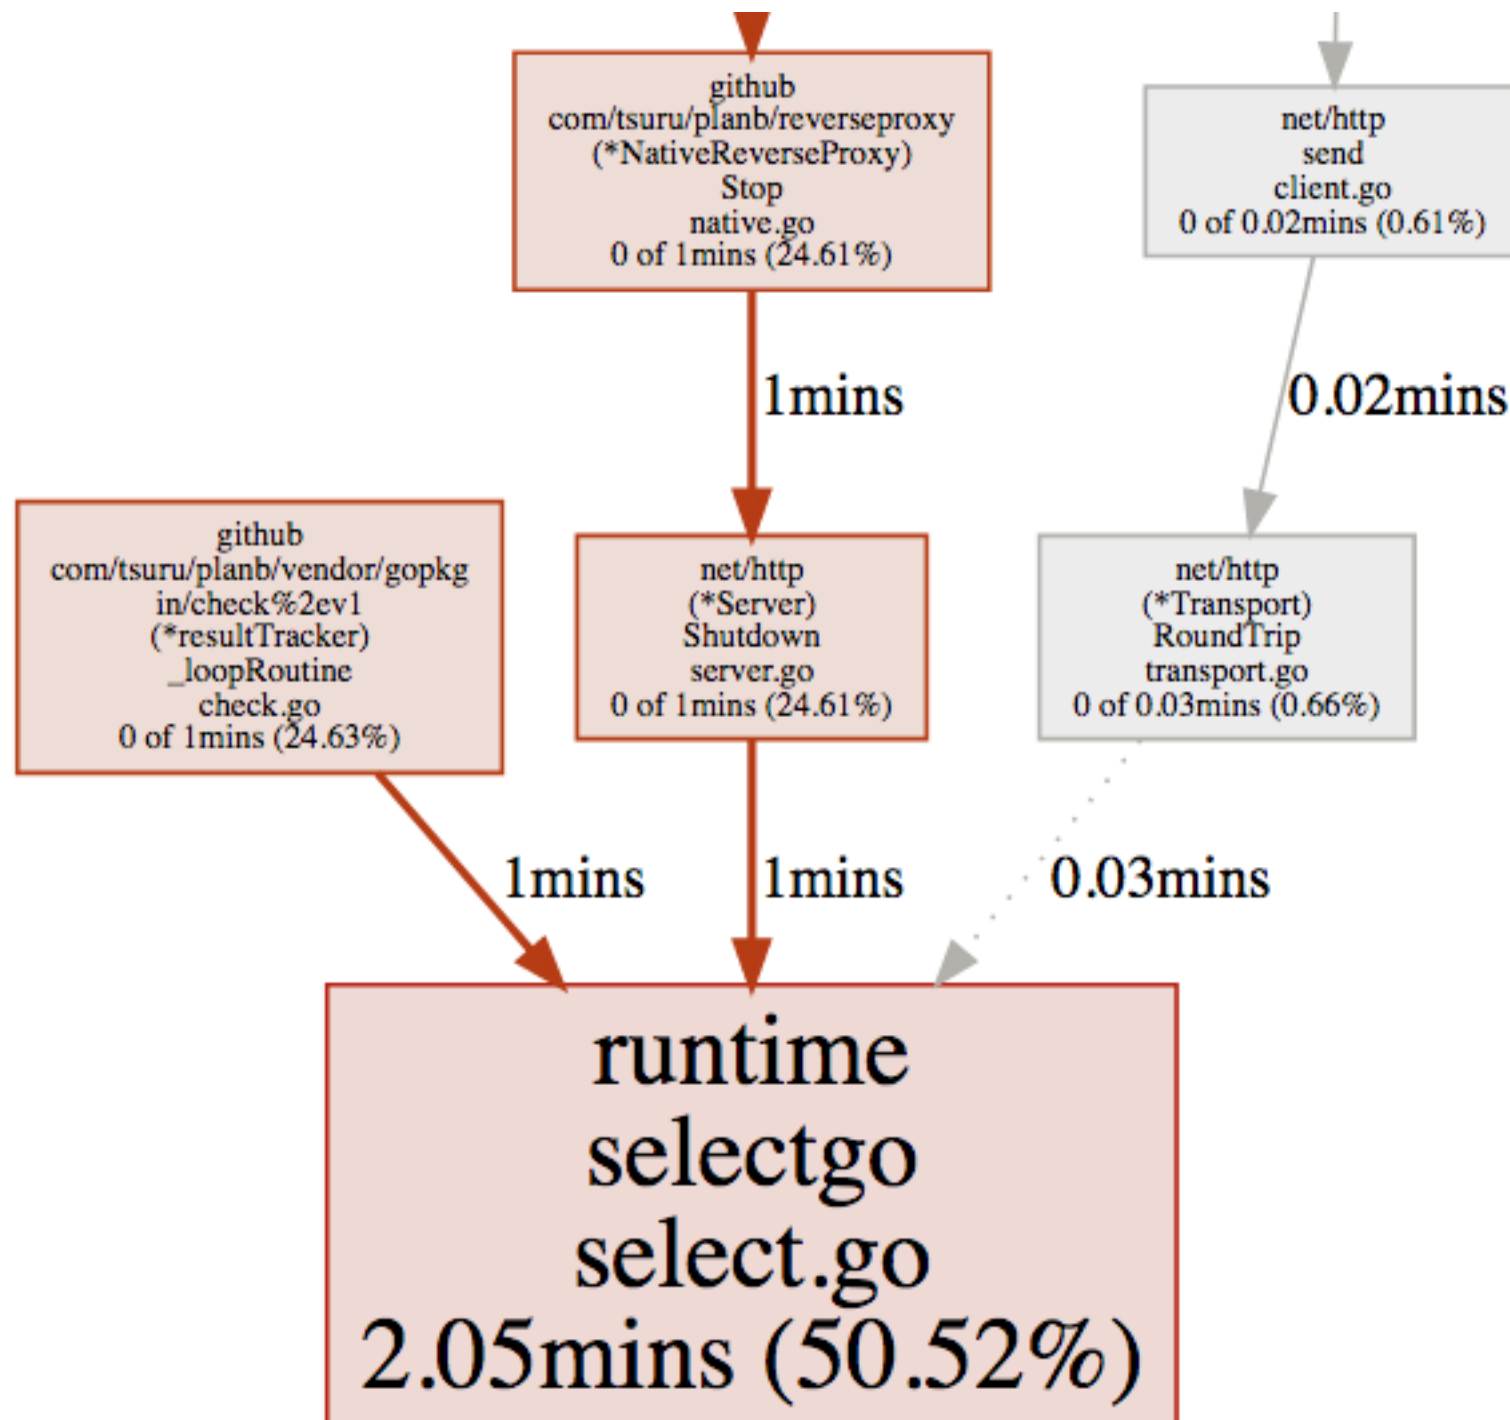
```

763 func (s *S) TestRoundTripStressWithTimeoutBackend(c *check.C) {
764     router := &noopRouter{dst: "http://127.0.0.1:23771"}
765     rp := s.factory()
766     err := rp.Initialize(ReverseProxyConfig{Router: router})
767     c.Assert(err, check.IsNil)
768     addr, listener := getFreeListener()
769     go rp.Listen(listener, nil)
770     defer rp.Stop()
771     defer listener.Close()
772     u := fmt.Sprintf("http://%s/", addr)
773     wg := sync.WaitGroup{}
774     nConnections := 50
775     for i := 0; i < nConnections; i++ {
776         wg.Add(1)
777         go func() {
778             defer wg.Done()
779             req, err := http.NewRequest("GET", u, nil)
780             c.Assert(err, check.IsNil)
781             req.Host = "myhost.com"
782             rsp, err := http.DefaultClient.Do(req)
783             c.Assert(err, check.IsNil)
784             defer rsp.Body.Close()
785             c.Assert(rsp.StatusCode, check.Equals, 503)
786         }()
787     }
788     done := make(chan bool)
789     go func() {
790         wg.Wait()
791         close(done)
792     }()
793     select {
794     case <-done:
795     case <-time.After(time.Minute):
796         c.Fatal("timeout out after 1 minute")
797     }
798 }

```

Let's trace it!

```
$ go test -trace=trace.out -check.v -check.f TestRoundTripStressWithTimeout
```



```
2487 func (srv *Server) Shutdown(ctx context.Context) error {
2488     atomic.AddInt32(&srv.inShutdown, 1)
2489     defer atomic.AddInt32(&srv.inShutdown, -1)
2490
2491     srv.mu.Lock()
2492     lnerr := srv.closeListenersLocked()
2493     srv.closeDoneChanLocked()
2494     for _, f := range srv.onShutdown {
2495         go f()
2496     }
2497     srv.mu.Unlock()
2498
2499     ticker := time.NewTicker(shutdownPollInterval)
2500     defer ticker.Stop()
2501     for {
2502         if srv.closeIdleConns() {
2503             return lnerr
2504         }
2505         select {
2506             case <-ctx.Done():
2507                 return ctx.Err()
2508             case <-ticker.C:
2509             }
2510     }
2511 }
```

net/http.(*conn).serve N=50

Goroutine	Total time, ns	Execution time, ns	Network wait time, ns	Sync block time, ns	Blocking syscall time, ns	Scheduler wait time, ns	GC sweeping time, ns	GC pause time, ns
432	60036424682	20108	60035899328	0	0	505246	0	7890133
49	14275321	568388	1062815	11064154	0	1579966	0	0
149	14227546	496880	658666	10361225	0	2710775	0	0
152	14210924	222502	2766080	7996838	0	3225504	0	0
147	14159859	528756	1350353	10095634	0	2185116	0	0

<https://github.com/golang/go/issues/21204>

Workaround

⚙	@@ -763,7 +763,7 @@ func (s *S) TestRoundTripStress(c *check.C) {	
763	763	func (s *S) TestRoundTripStressWithTimeoutBackend(c *check.C) {
764	764	router := &noopRouter{dst: "http://127.0.0.1:23771"}
765	765	rp := s.factory()
766	-	err := rp.Initialize(ReverseProxyConfig{Router: router})
766	+	err := rp.Initialize(ReverseProxyConfig{Router: router, ReadTimeout: time.Second})
767	767	c.Assert(err, check.IsNil)
768	768	addr, listener := getFreeListener()
769	769	go rp.Listen(listener, nil)
⚙	@@ -1041,7 +1041,7 @@ func baseBenchmarkServeHTTP(rp ReverseProxy, b *testing.B) {	

```
$ go test -check.v -check.f TestRoundTripStressWithTimeout
PASS: reverseproxy_test.go:763: S.TestRoundTripStressWithTimeoutBackend 1.013s
PASS: reverseproxy_test.go:763: S.TestRoundTripStressWithTimeoutBackend 0.013s
OK: 2 passed
PASS
ok      github.com/tsuru/planb/reverseproxy    1.055s
```

<https://github.com/tsuru/planb/pull/35>

Conclusions

- Go execution tracer helps understand concurrency
 - Complements other tools (eg memory/cpu profile)
- Not much documentation available
 - Opportunity for contributions

Reference

- Using the go tracer to speed up fractal rendering
 - Just for func #22
- Go Execution Tracer (Design Doc)

Obrigado!

Slides and notes in english
will be available at
<https://andrestc.com>